

Task - 1

```
import pandas as pd
import numpy as np

data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
data = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]
```

```
raw_df.head()
```

↗

	0	1	2	3	4	5	6	7	8	9	10
0	0.00632	18.00	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3
1	396.90000	4.98	24.00	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2	0.02731	0.00	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8
3	396.90000	9.14	21.60	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
4	0.02729	0.00	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8

```
# import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings("ignore")
```

```
df = pd.read_csv("https://mjnvsai.github.io/Machine_Learning/heart.csv")

# The shape of the data
print("The shape of the dataset is : ", df.shape)

# Preview of the first 5 rows of the data
df.head()
```

The shape of the dataset is : (303, 14)

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	thall	output
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

```
df.describe()
```

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh		
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	0.528053	149.646865	0.000000	0.000000
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	0.525860	22.905161	0.000000	0.000000
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.000000	0.000000	0.000000
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000	133.500000	0.000000	0.000000
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000	153.000000	0.000000	0.000000
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	1.000000	166.000000	1.000000	1.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000	202.000000	1.000000	1.000000

```
# Checking the number of unique values in each column

dict = {}
for i in list(df.columns):
    dict[i] = df[i].value_counts().shape[0]
```

```
pd.DataFrame(dict,index=["unique count"]).transpose()
```

unique count	
age	41
sex	2
cp	4
trtbps	49
chol	152
fbs	2
restecg	3
thalachh	91
exng	2
oldpeak	40
slp	3
caa	5
thall	4
output	2

```
# Separating the columns in categorical and continuous

cat_cols = ['sex','exng','caa','cp','fbs','restecg','slp','thall']
con_cols = ["age","trtbps","chol","thalachh","oldpeak"]
target_col = ["output"]
print("The categorial cols are : ", cat_cols)
print("The continuous cols are : ", con_cols)
print("The target variable is : ", target_col)

The categorial cols are :  ['sex', 'exng', 'caa', 'cp', 'fbs', 'restecg', 'slp', 'thall']
The continuous cols are :  ['age', 'trtbps', 'chol', 'thalachh', 'oldpeak']
The target variable is :  ['output']
```

```
# Summary statistics

df[con_cols].describe().transpose()
```

	count	mean	std	min	25%	50%	75%	max
age	303.0	54.366337	9.082101	29.0	47.5	55.0	61.0	77.0
trtbps	303.0	131.623762	17.538143	94.0	120.0	130.0	140.0	200.0
chol	303.0	246.264026	51.830751	126.0	211.0	240.0	274.5	564.0
thalachh	303.0	149.646865	22.905161	71.0	133.5	153.0	166.0	202.0
oldpeak	303.0	1.039604	1.161075	0.0	0.0	0.8	1.6	6.2

```
# Missing values

df.isnull().sum()
```

age	0
sex	0
cp	0
trtbps	0
chol	0
fbs	0
restecg	0
thalachh	0
exng	0
oldpeak	0
slp	0
caa	0
thall	0
output	0
dtype: int64	

```
# Scaling
from sklearn.preprocessing import RobustScaler

# Train Test Split
```

```

from sklearn.model_selection import train_test_split

# Models
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import GradientBoostingClassifier

# Metrics
from sklearn.metrics import accuracy_score, classification_report, roc_curve

```

```

# Scaling and Encoding features

# creating a copy of df
df1 = df

# define the columns to be encoded and scaled
cat_cols = ['sex', 'exng', 'caa', 'cp', 'fbs', 'restecg', 'slp', 'thall']
con_cols = ["age", "trtbps", "chol", "thalachh", "oldpeak"]

# encoding the categorical columns
df1 = pd.get_dummies(df1, columns = cat_cols, drop_first = True)

# defining the features and target
X = df1.drop(['output'], axis=1)
y = df1[['output']]

# instantiating the scaler
scaler = RobustScaler()

# scaling the continuous featuree
X[con_cols] = scaler.fit_transform(X[con_cols])
print("The first 5 rows of X are")
X.head()

```

The first 5 rows of X are

	age	trtbps	chol	thalachh	oldpeak	sex_1	exng_1	caa_1	caa_2	caa_3	...	cp_2	cp_3	1
0	0.592593	0.75	-0.110236	-0.092308	0.9375	1	0	0	0	0	...	0	1	
1	-1.333333	0.00	0.157480	1.046154	1.6875	1	0	0	0	0	...	1	0	
2	-1.037037	0.00	-0.566929	0.584615	0.3750	0	0	0	0	0	...	0	0	
3	0.074074	-0.50	-0.062992	0.769231	0.0000	1	0	0	0	0	...	0	0	
4	0.148148	-0.50	1.795276	0.307692	-0.1250	0	1	0	0	0	...	0	0	

5 rows × 22 columns

```

# Dimensionality Reduction

from sklearn.decomposition import PCA

pca = PCA(n_components=22)
X = pca.fit_transform(X)

```

```

# Train and test split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)
print("The shape of X_train is: ", X_train.shape)
print("The shape of X_test is: ", X_test.shape)
print("The shape of y_train is: ", y_train.shape)
print("The shape of y_test is: ", y_test.shape)

```

```

The shape of X_train is: (242, 22)
The shape of X_test is: (61, 22)
The shape of y_train is: (242, 1)
The shape of y_test is: (61, 1)

```

```

# instantiating the object and fitting
clf = SVC(kernel='linear', C=1, random_state=42).fit(X_train, y_train)

# predicting the values
y_pred = clf.predict(X_test)

# printing the test accuracy
print("The test accuracy score of SVM is ", accuracy_score(y_test, y_pred))
print("Classification Report: \n", classification_report(y_test, y_pred, digits=4))

```

The test accuracy score of SVM is 0.8688524590163934
 Classification Report:

	precision	recall	f1-score	support
0	0.8387	0.8966	0.8667	29
1	0.9000	0.8438	0.8710	32
accuracy			0.8689	61
macro avg	0.8694	0.8702	0.8688	61
weighted avg	0.8709	0.8689	0.8689	61

```
# define logistic regression impl method

class LogisticRegression:

    def __init__(self, alpha=0.01, iters=1000, verbose = False):
        self.alpha = alpha
        self.iters = iters
        self.theta = None
        self.verbose = verbose

    # Logistic Function
    def sigmoid(self, z):
        return 1 / (1 + np.exp(-z))

    # Cost Function
    def cost(self, h, y):
        return (-y * np.log(h) - (1 - y) * np.log(1 - h)).mean()

    #Gradient Computation
    def gradient(self, X, h, y):
        return np.dot(X.T, (h - y)) / y.shape[0]

    # Parameter Update
    def update_theta(self, gradient, lr):
        return self.theta - (gradient * self.alpha)

    def fit(self, X, y):
        # Initial theta values
        np.random.seed(999)
        self.theta = np.random.randn(X.shape[1])
        cost_array = np.zeros(self.iters)

        for i in range(self.iters):
            h = self.sigmoid(np.dot(X, self.theta))
            cost_num = self.cost(h, y)
            cost_array[i] = cost_num
            gradient = self.gradient(X, h, y)
            self.theta = self.update_theta(gradient, self.alpha)

        # Print training History
        if(self.verbose):
            if(self.iters<=1000):
                if(i%100==0):
                    print(f"Iter {i}: cost: {cost_array[i]}")
            elif(self.iters<=10000 and self.iters>1000):
                if(i%1000==0):
                    print(f"Iter {i}: cost: {cost_array[i]}")
            else:
                if(i%10000==0):
                    print(f"Iter {i}: cost: {cost_array[i]}")

        return cost_array

    def predict(self, X):
        h = self.sigmoid(np.dot(X, self.theta))
        preds = np.where(h>=0.5, 1, 0)
        return np.array(preds)

    def predict_proba(self, X):
        h = self.sigmoid(np.dot(X, self.theta))
        return np.array(h)

    def plotChart(self, cost_num):
        fig, ax = plt.subplots()
        ax.plot(np.arange(self.iters), cost_num, 'r')
        ax.set_xlabel('Iterations')
        ax.set_ylabel('Cost')
        ax.set_title('Error vs Iterations')
        plt.show()
```

```
# hyper params
lr = 0.005
iters = 100000

log_reg = LogisticRegression(lr, iters, True)
costs = log_reg.fit(X_train, y_train.values.reshape(242,))
log_reg.plotChart(costs)
y_pred = log_reg.predict(X_test)

# printing the test accuracy
print("The test accuracy score of Logistric Regression is ", accuracy_score(y_test, y_pred))
print("Classification Report: \n", classification_report(y_test, y_pred, digits=4))

# calculating the probabilities
y_pred_prob = log_reg.predict_proba(X_test)

# instantiating the roc_cruve
fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)

# plotting the curve
plt.plot([0,1],[0,1], "k--", 'r+')
plt.plot(fpr, tpr, label='Logistic Regression')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Logistric Regression ROC Curve")
plt.show()
```

▼ Task - 2

```

Iter 0: COST: 0.9630/90900890911
import pandas as pd
import numpy as np
Iter 40000: COST: 0.304/201/33044102
df = pd.read_table('https://mjnvai.github.io/Machine_Learning/fruit.txt')
Iter 10000: COST: 0.2951435180224018
df.sample(10)

```

	fruit_label	fruit_name	fruit_subtype	mass	width	height	color_score
50	4	lemon	unknown	130	6.0	8.2	0.71
46	4	lemon	spanish_belsan	216	7.3	10.2	0.71
41	3	orange	turkey_navel	180	7.6	8.2	0.79
15	1	apple	golden_delicious	156	7.7	7.1	0.69
32	3	orange	selected_seconds	164	7.2	7.0	0.80
56	4	lemon	unknown	116	5.9	8.1	0.73
2	1	apple	granny_smith	176	7.4	7.2	0.60
10	1	apple	braeburn	166	6.9	7.3	0.93
47	4	lemon	spanish_belsan	196	7.3	9.7	0.72
13	1	apple	golden_delicious	164	7.3	7.7	0.70

```

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 59 entries, 0 to 58
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   fruit_label      59 non-null    int64
1   fruit_name       59 non-null    object
2   fruit_subtype    59 non-null    object
3   mass             59 non-null    int64
4   width            59 non-null    float64
5   height           59 non-null    float64
6   color_score      59 non-null    float64
dtypes: float64(3), int64(2), object(2)
memory usage: 3.4+ KB

```

```

X = df.iloc[:,3:7].values
y = df.iloc[:,1].values
print('Total Features are', X.shape)
print('Total Labels are',y.shape)

```

```

Total Features are (59, 4)
Total Labels are (59,)

```

```

from sklearn.model_selection import train_test_split
# Splitting the dataset into the Training set and Test set from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,random_state=10)

```

```
print(X_train.shape,X_test.shape)
```

```
(47, 4) (12, 4)
```

```

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.metrics import classification_report

```

```

logreg = LogisticRegression() # Create an instance of Logistic Regression Classifier and fit the data.
model = logreg.fit(X_train,y_train)

```

```
y_pred=model.predict(X_test)
```

```

print('Actual',"--","Predicted")
for i in range(len(y_pred)):
    print(y_test[i],"--",y_pred[i])

```

```

Actual -- Predicted
orange -- orange
mandarin -- mandarin

```

```

orange -- apple
orange -- orange
apple -- apple
apple -- apple
lemon -- lemon
apple -- apple
apple -- apple
orange -- orange
orange -- lemon
orange -- apple

```

```

print("Accuracy: ",accuracy_score(y_test, y_pred)*100)

print("\n***Confusion Matrix***\n",confusion_matrix(y_test, y_pred))

print("\n***Classification report***\n",classification_report(y_test, y_pred))

```

```

Accuracy:  75.0

***Confusion Matrix***
[[4 0 0 0]
 [0 1 0 0]
 [0 0 1 0]
 [2 1 0 3]]

***Classification report***
              precision    recall  f1-score   support

   apple         0.67         1.00         0.80         4
   lemon         0.50         1.00         0.67         1
  mandarin         1.00         1.00         1.00         1
   orange         1.00         0.50         0.67         6

   accuracy              0.75              12
  macro avg         0.79         0.88         0.78         12
 weighted avg         0.85         0.75         0.74         12

```

```

#import KNClassifier
from sklearn.neighbors import KNeighborsClassifier

#Create an Instance for KN Classifier
neigh = KNeighborsClassifier(n_neighbors=3)

# Perform Training
model = neigh.fit(X_train, y_train)

```

```

print('Actual',"--","Predicted")
for i in range(len(y_pred)):
    print(y_test[i],"--",y_pred[i])

```

```

Actual -- Predicted
orange -- orange
mandarin -- mandarin
orange -- apple
orange -- orange
apple -- apple
apple -- apple
lemon -- lemon
apple -- apple
apple -- apple
orange -- orange
orange -- lemon
orange -- apple

```

```

#making predictions
y_pred=model.predict(X_test)

print("Accuracy: ",accuracy_score(y_test, y_pred)*100)

print("\n***Confusion Matrix***\n",confusion_matrix(y_test, y_pred))

print("\n***Classification report***\n",classification_report(y_test, y_pred))

```

```

Accuracy:  75.0

***Confusion Matrix***
[[4 0 0 0]
 [0 1 0 0]
 [0 0 1 0]
 [1 2 0 3]]

***Classification report***
              precision    recall  f1-score   support

   apple         0.67         1.00         0.80         4
   lemon         0.50         1.00         0.67         1
  mandarin         1.00         1.00         1.00         1
   orange         1.00         0.50         0.67         6

   accuracy              0.75              12
  macro avg         0.79         0.88         0.78         12
 weighted avg         0.85         0.75         0.74         12

```

apple	0.80	1.00	0.89	4
lemon	0.33	1.00	0.50	1
mandarin	1.00	1.00	1.00	1
orange	1.00	0.50	0.67	6
accuracy			0.75	12
macro avg	0.78	0.88	0.76	12
weighted avg	0.88	0.75	0.75	12

```
from sklearn.tree import DecisionTreeClassifier

# Create Model
decision_tree = DecisionTreeClassifier(criterion='gini')
# Perform the training
model = decision_tree.fit(X_train, y_train)
```

```
#making predictions
y_pred=model.predict(X_test)
```

```
print('Actual',"--","Predicted")
print("-----")
for i in range(len(y_pred)):
    print(y_test[i],"--",y_pred[i])
```

```
Actual -- Predicted
-----
orange -- orange
mandarin -- mandarin
orange -- orange
orange -- orange
apple -- apple
apple -- orange
lemon -- lemon
apple -- apple
apple -- apple
orange -- orange
orange -- orange
orange -- orange
```

```
print("Accuracy: ",accuracy_score(y_test, y_pred)*100)

print("\n***Confusion Matrix***\n",confusion_matrix(y_test, y_pred))

print("\n***Classification report***\n",classification_report(y_test, y_pred))
```

```
Accuracy: 91.66666666666666
```

```
***Confusion Matrix***
[[3 0 0 1]
 [0 1 0 0]
 [0 0 1 0]
 [0 0 0 6]]
```

```
***Classification report***
              precision    recall  f1-score   support

   apple         1.00        0.75        0.86         4
   lemon         1.00        1.00        1.00         1
  mandarin         1.00        1.00        1.00         1
    orange         0.86        1.00        0.92         6

   accuracy              0.92         12
  macro avg              0.96        0.94        0.95         12
 weighted avg              0.93        0.92        0.91         12
```



```
#import SVM Classifier
from sklearn.svm import SVC

#Create an Instance for KN Classifier
svm_l = SVC(kernel = 'linear')
# Perform Training
model = svm_l.fit(X_train, y_train)
#making predictions
y_pred=model.predict(X_test)

print("Accuracy: ",accuracy_score(y_test, y_pred)*100)

print("\n***Confusion Matrix***\n",confusion_matrix(y_test, y_pred))

print("\n***Classification report***\n",classification_report(y_test, y_pred))
```

Accuracy: 83.33333333333334

Confusion Matrix

```
[[3 0 0 1]
 [0 1 0 0]
 [0 0 1 0]
 [0 1 0 5]]
```

Classification report

	precision	recall	f1-score	support
apple	1.00	0.75	0.86	4
lemon	0.50	1.00	0.67	1
mandarin	1.00	1.00	1.00	1
orange	0.83	0.83	0.83	6
accuracy			0.83	12
macro avg	0.83	0.90	0.84	12
weighted avg	0.88	0.83	0.84	12

```
#import SVM Classifier
from sklearn.svm import SVC

#Create an Instance for KN Classifier
svm_r = SVC(kernel = 'rbf')
# Perform Training
model = svm_r.fit(X_train, y_train)
#making predictions
y_pred=model.predict(X_test)

print("Accuracy: ",accuracy_score(y_test, y_pred)*100)

print("\n***Confusion Matrix***\n",confusion_matrix(y_test, y_pred))

print("\n***Classification report***\n",classification_report(y_test, y_pred))
```

Accuracy: 41.66666666666667

Confusion Matrix

```
[[4 0 0 0]
 [1 0 0 0]
 [0 1 0 0]
 [4 1 0 1]]
```

Classification report

	precision	recall	f1-score	support
apple	0.44	1.00	0.62	4
lemon	0.00	0.00	0.00	1
mandarin	0.00	0.00	0.00	1
orange	1.00	0.17	0.29	6
accuracy			0.42	12
macro avg	0.36	0.29	0.23	12
weighted avg	0.65	0.42	0.35	12

```
#import SVM Classifier
from sklearn.svm import SVC

#Create an Instance for KN Classifier
svm_p = SVC(kernel = 'poly',degree=2)
# Perform Training
model = svm_p.fit(X_train, y_train)
#making predictions
y_pred=model.predict(X_test)
```

```
print("Accuracy: ",accuracy_score(y_test, y_pred)*100)

print("\n***Confusion Matrix***\n",confusion_matrix(y_test, y_pred))

print("\n***Classification report***\n",classification_report(y_test, y_pred))
```

Accuracy: 41.66666666666667

Confusion Matrix

```
[[4 0 0 0]
 [1 0 0 0]
 [0 1 0 0]
 [3 2 0 1]]
```

Classification report

	precision	recall	f1-score	support
apple	0.50	1.00	0.67	4
lemon	0.00	0.00	0.00	1
mandarin	0.00	0.00	0.00	1
orange	1.00	0.17	0.29	6
accuracy			0.42	12
macro avg	0.38	0.29	0.24	12
weighted avg	0.67	0.42	0.37	12

